

Human-like autonomy emerges from self-play and a pinch of human data

Daphne Cornelisse¹ Julian Hunt² Zixu Zhang³ Waël Doulazmi^{4,5} Kevin Joseph²

Jaime Fernández Fisac³ Eugene Vinitsky¹

¹NYU Tandon School of Engineering ²NYU Courant ³Princeton University
⁴Centre for Robotics, Mines Paris ⁵Valeo

Abstract: Self-play reinforcement learning has recently taken the wheel in autonomous driving. This approach uses cheap, large-scale simulation to substitute expensive, large-scale human demonstration data. While cost-effective, agents trained through self-play often behave unpredictably around unseen policies and fail to coordinate with human drivers due to incompatible driving styles. Previous works attempt to mitigate this behavioral incompatibility through reward engineering, which is labor-intensive and brittle. In this work, we build on the best of both approaches and learn driving policies by training on over 30 years of self-play simulation against a minimal reward function and just 30 minutes of human demonstrations. Our agents achieve near-zero-shot coordination with real human driving trajectories using $125,000 \times$ less human data than imitation learning baselines, at 1/25 the number of GPU hours. The solution space of unconstrained driving policies is large, but this small pinch of human data is sufficient to act as a regularizer towards ones that fit the road. Together, this suggests that human-compatible agents can be trained with far less human demonstration data than is currently assumed.

Keywords: Zero-Shot Coordination, Self-play Reinforcement Learning, Imitation Learning, Autonomous Driving

1 Introduction

Self-play reinforcement learning (RL) has produced superhuman agents in strategic games [1, 2, 3] and, more recently, has shown promise in real-world domains including autonomous driving [4, 5, 6, 7] and robotic manipulation [8]. The approach elegantly sidesteps a central difficulty in multi-agent learning - how to model the opponent - through the following idea: the *agent’s opponent is a copy of itself*. The appeal here is that as the agent improves, so does its co-player. This gives rise to an automatically evolving curriculum [9] that takes behavior from random play to optimal behavior entirely through synthetic simulated experience.

In zero-sum games, this mechanism with a sparse measure for success (e.g., +1 when winning a game of chess) is enough to produce strong play against arbitrary opponents. Many real-world settings, however, are not zero-sum. Driving, for instance, is a mixed-motive game: each agent has *individual objectives* (reaching a destination safely) but must also *coordinate* with other road users by adhering to shared norms, expectations, and conventions. Self-play RL with only a high-level objective for success provides no guarantees of such alignment; agents may converge to effective but “alien” strategies that are incompatible with human partners [10]. Concretely, an agent trained to “reach a destination safely,” may very well learn to do so in reverse, sideways, or on the wrong side of the road if such constraints are not specified in the reward.




Figure 1: **TODO** Regularized self-play RL reduces the human-coordination gap with a pinch of human demonstration data. Left: score on a human-replay evaluation (proxy for human-compatibility) versus human demonstration data per method. Adding 30 min of human demonstrations as an RL anchor (star, ours) increases performance well above the unregularized self-play baseline (square) and SMART (circle), an imitation-learning model trained on the full Waymo dataset. Right: total training transitions per method. Our method is trained on substantially more data than SMART, which explains the performance improvement. However, in contrast to SMART, we train almost entirely on cheap synthetic data, whereas SMART trains on 45M human demonstration transitions (52 days).

Previous works have addressed this misalignment in two ways. One line of work involves *manual reward engineering*, where reward terms are added iteratively until the desired behavior and conventions emerge [5, 11]. While effective, this strategy is labor-intensive by nature, domain-specific, and brittle since it is not trivial to figure out what reward will produce the desired human-like behavior [12]. A good example is GIGAFLow [5], which required nine individually tuned reward terms and several other domain randomization tricks to produce naturalistic driving behavior. On the other side of the spectrum we have *Imitation learning* [13, 14, 15]. IL directly imitates human demonstrations to avoid having to define a reward function altogether. However, robustness requires broad behavioral coverage, so these approaches typically require large quantities of human demonstrations to work well [16].

We take a different approach, grounded in a practical observation about the changing cost structure of experience generation. Modern RL frameworks and simulation infrastructure can generate between 300K and 20M environment steps *per second* on a single consumer-grade GPU [17, 18], making synthetic experience effectively limitless. Human demonstrations, by contrast, require manual collection and remain difficult to scale. This suggests a natural role for human demonstration data in games with the aforementioned properties: not as the primary source of training signal, but as a lightweight anchor that steers the policy away from effective yet behaviorally alien strategies. Indeed, regularizing self-play RL toward such an anchor has shown promise in producing human-compatible agents in Diplomacy [19, 20] and driving [21, 22, 7], yet *how much* demonstration data is required to reach human compatibility remains, to our knowledge, unexamined.

We present the first systematic study of this question. We regularize self-play PPO toward a behavioral cloning anchor trained on varying amounts of human demonstrations from the Waymo Open Motion Dataset [23, WOMD], and find that the answer is: *surprisingly little*. Human-compatible driving agents can be learned from as little as **30 minutes to 3 hours of human demonstrations**, roughly **0.25%** of the human demonstration data used to train state-of-the-art imitation learning models, while using 1/25 of the GPU hours times faster. More broadly, the present results suggest that **wherever simulation is cheap, collecting a small amount of human demonstrations may suffice** to learn agents that are compatible with human play.

Our contributions are:

- We show that regularizing self-play RL toward a behavioral cloning anchor achieves improved coordination with human drivers where pure self-play fails, without reward engineering and with 30 minutes to 3 hours of human demonstration data when combined with 30 years of self-play generated synthetic data.
- We present the first systematic study of human data scaling under regularized self-play RL in the application domain of autonomous driving.
- We open-source the full codebase; each training run is reproducible in seven hours on a single consumer-class GPU.

2 Related Work

Imitation learning for autonomous driving. The generation of driving policies is a fundamental challenge across end-to-end autonomous driving [24, 25, 26, 27], multi-agent trajectory prediction [28], and reactive traffic simulation [29, 30]. Driven by the widespread availability of large-scale human driving datasets [31, 23, 32], imitation learning has become the dominant approach across all these domains [33]. Under this imitation learning paradigm, a broad spectrum of methodologies has emerged to fit models to historical data, ranging from marginal [34, 35, 36] and joint [37, 38, 39, 40] forecasting to autoregressive sequence modeling of tokenized trajectories [41, 15, 42] and continuous distribution learning via diffusion and promptable world models [43, 44, 45, 46, 47]. While these generative approaches yield diverse open-loop behaviors, they are fundamentally constrained by the scale of human data required and frequently suffer from compounding covariate shift in closed-loop deployment [16]. To mitigate these shifts, recent hybrid approaches integrate reinforcement learning [48, 49, 50], yet they typically still rely on extensive human demonstrations as their primary optimization signal. Our approach systematically inverts this balance: rather than depending on human driving logs as the core supervisor, we utilize synthetic, multi-agent RL self-play as the primary engine for discovering robust interactive behaviors, retaining a remarkably small human dataset strictly as a behavioral anchor to ensure conformity to realistic traffic norms.

Self-play reinforcement learning in games. Self-play reinforcement learning has produced super-human agents in games from Go and Chess [1, 51] to StarCraft II [52] and Stratego [2], all without human data. Superhuman play is not the same as human-compatible play, however. Many games admit multiple equilibria, and self-play need not converge to equilibria that are compatible with human partners [10, 53]. The failure has been shown in cooperative games such as Hanabi [54] and Diplomacy [10], where self-play agents develop internally consistent conventions that transfer poorly to human partners. The cause is reward underspecification: when the reward is defined as a score to maximize, there are often many ways to achieve it. In other words, the solution space is large. Previous work attempts to resolve this by sharing the reward by hand [5, 11]. For instance, GIGAFLOW [5] demonstrates that reward engineering and domain randomization can produce naturalistic behavior at scale, at the cost of nine individually tuned reward terms. We avoid reward engineering entirely. A small amount of human data serves as a behavioral anchor, and self-play does the rest. This reduces a labor-intensive design problem to a one-hour data collection.

Human-regularized self-play reinforcement learning and search. One alternative to reward engineering is to regularize self-play toward a human anchor policy. This idea has been explored in Diplomacy, where KL regularization toward a human prior during both search and learning produced agents that coordinate more effectively with human partners [19, 20]. Jacob et al. [55] study KL-regularized search more broadly and show that it recovers human-like play across several games. In autonomous driving, Cornelisse and Vinitzky [21] show that regularizing PPO toward a behavioral cloning anchor produces more human-compatible driving agents on a small scale. More recently, Chang et al. [7] show that KL-regularized self-play can train human-like driving policies from scratch using SMART [42] as the behavioral anchor, a large tokenized model trained on the full 500,000-scenario Waymo dataset. [22] show that regularized self-play RL effectively adapts policies to new cities, without using human demonstrations from that city. In sum, the effectiveness of




Figure 2: **Evaluation settings.** Self-play (left) and human-replay (center, right). Red arrows mark collisions. Rectangles are vehicles; squares are pedestrians. In human-replay, some collisions are effectively unavoidable: replay agents follow their logged trajectories and can drive into the controlled SDC from behind. We therefore distinguish between *collisions* (any contact) and *at-fault collisions* (contact caused by the controlled agent, following NAVSIM [58]).

regularized self-play RL has been shown. What remains unexamined is the scaling law that governs it: how performance depends on the amount and quality of the human anchor data. The question matters now in a way it did not before. Simulation throughput has grown by orders of magnitude in recent years, shifting the bottleneck from synthetic experience to human data.

3 Method

3.1 Evaluation and Metrics

Evaluation. A human-compatible agent should *blend in* with human drivers. Since on-road deployment is out of scope, we approximate the human population with logged human trajectories and evaluate in three settings, illustrated in Figure 2:

- **Self-play.** All agents are controlled by the trained policy.
- **Human-replay.** Only the self-driving car (SDC) is controlled by the policy; all other agents follow their logged trajectories.
- **IDM.** The SDC is controlled by the policy; all other agents follow the Intelligent Driver Model [56], tracking a precomputed lane-center path for lateral control and using IDM longitudinal accelerations to maintain a safe gap to the lead vehicle, as in [57].

A human-compatible agent should score well across all three: it should reach its goal without collisions or off-road events whether its co-players are copies of itself, reactive rule-based drivers, or human drivers it has never trained against. The three settings probe complementary failure modes. Human-replay is the strictest test: because logged trajectories cannot react, the SDC alone is responsible for avoiding collisions, so the setting stress-tests whether the policy has internalized human driving conventions rather than coordinating on idiosyncratic self-play conventions. IDM sits between self-play and human-replay in difficulty; co-players are no longer the policy itself (so self-play-specific conventions do not help), but they are reactive, so the SDC is not solely responsible for averting every collision.

Metrics and ceiling. Our primary metric is the **score**: an agent scores 1 if it completes the task of driving to a goal destination without colliding or going off-road, and 0 otherwise. To diagnose failure modes, we separately report **collision rate**, **at-fault collision rate**, **off-road rate**, and **route progress**. An ideal agent should score well with its own population as well as the human-replay population. Score-based metrics capture whether agents complete their task safely, but not whether their behavior looks human. We therefore also report **distributional realism** using the Waymo Open assess whether agents complete their tasks safely, but not whether their behavior appears of simulated

and logged trajectories across kinematic, interactive, and map-based features. All metrics are reported on held-out test scenarios unless stated otherwise; full definitions are given in Appendix D.2.

3.2 Simulation Environment

We train policies $\pi(a_t^i | o_t^i)$ in PufferDrive 2.0 [18], a batched driving simulator that reaches 300k SPS on a single consumer-grade GPU. Environments are initialized from the Waymo Open Motion Dataset [23]: each 9-second scenario provides a roadgraph, a variable set of agents (cars, cyclists, pedestrians) up to $N = 32$, and per-agent initial poses and goals drawn from the logged trajectories. Each agent receives a partial, decentralized, ego-frame observation consisting of its own state, the $N - 1$ closest neighbors within 50 m, and up to 128 nearby road segments.

We avoid reward shaping and use a sparse reward driven only by high-level outcomes: +1 for reaching the goal position, -1 for collision or going off-road, and 0 otherwise. This way, human-like behavior emerges from regularization rather than from a hand-tuned reward. Episodes are truncated once all agents reach their destinations. Full details are given in Appendix A.

3.3 Regularized Self-Play Reinforcement Learning

Our method has two stages: we first train a behavioral cloning (BC) anchor on human demonstrations, then use it as a regularizer during self-play RL. Note that the RL policy is learned from scratch rather than fine-tuned from the anchor.

Step 1: Train the anchor policy. To study how the amount of human demonstration data affects downstream performance, we train a set of anchor policies on nested subsets of the demonstrations. Given the full dataset $\mathcal{D} = \{(o_t^i, a_t^i)\}_{i=1}^{T \cdot K}$ of observation-action pairs from K scenarios of T timesteps each, we sample subsets $\mathcal{D}_n \subset \mathcal{D}$ corresponding to n scenarios, yielding roughly {10 min, 30 min, 3 h, 30 h} of human driving data. Each anchor τ_n is trained by minimizing the negative log-likelihood of the logged actions:

$$\tau_n = \arg \min_{\tau \in \mathcal{T}} \sum_{(o_t^i, a_t^i) \in \mathcal{D}_n} -\log \tau(a_t^i | o_t^i), \quad (1)$$

Figure 10 illustrates the data collection process in three scenarios. Since we have enough scenarios, we only use the SDC trajectory from each scenario, as it is typically the highest-quality trajectory in the scene. Each τ_n is then frozen and used as the regularizer during the subsequent self-play stage. See Appendix A.4 for the full details.

Step 2: Regularized self-play RL. We train a new policy π_θ from scratch using PPO [59] in self-play, with each anchor τ_n acting as a behavioral regularizer. Concretely, we add a KL penalty from the anchor to the policy to the standard PPO objective:

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{PPO}}(\theta) + \lambda \mathbb{E}_{o \sim \rho_{\pi_\theta}} [D_{\text{KL}}(\tau_n(\cdot | o) \parallel \pi_\theta(\cdot | o))], \quad (2)$$

where ρ_{π_θ} is the state distribution induced by the on-policy rollouts and $\lambda \geq 0$ controls the regularization strength. The KL term pulls π_θ toward the anchor’s action distribution on states the policy actually visits, rather than on the distribution of states in $\mathcal{D}_n$. We use scenario metadata to initialize worlds during self-play (Appendix A.1). The full training objective and hyperparameters are given in Appendix B.

4 Experiments

This section summarizes the key results. Additional details and experiments are reported in the appendices, including details on datasets, network architectures, hyperparameters, and environment settings. The appendices also report ablations.

The sections are structured to answer the following questions:

1. *Scaling laws for human demonstration data*: How much human demonstration data is needed for strong performance in both the self-play and human-replay settings? How do these scaling laws compare to those of the state-of-the-art imitation learning baseline? (Section 4.1)
2. *The role of metadata and scenario diversity*: Driving datasets such as WOMB and NuPlan provide scenario metadata—road graphs and initial agent positions—that grounds simulation at a fraction of the cost of collecting human demonstrations. How does the number of training scenarios (maps) used for self-play influence agent performance? (Section 4.2)
3. *Behavioral analysis*: Human-likeness in driving is not easily captured by any single metric. We conduct a full behavioral analysis comparing regularized self-play agents against their unregularized counterparts and the imitation learning baseline. (Section 4.3). Videos to show the difference in agent behavior are provided on an anonymous site: <https://sites.google.com/view/anonymous-human-like-autonomy>

4.1 Scaling Laws for Human Demonstrations

How much human demonstration data does regularized self-play need, and how does this compare to pure imitation learning? The second question matters because any apparent data efficiency on our side could simply reflect the homogeneity of the Waymo Open Dataset rather than an actual property of the method.

To answer this, we benchmark against SMART-tiny-CLSFT [42, 50] (—●—), the state-of-the-art imitation-learning approach in this field. We train SMART on the same nested subsets of human demonstrations as regularized self-play (—●—). The two pipelines differ only in the second stage, where SMART-tiny-CLSFT applies CAT-K closed-loop fine-tuning [50] (see Appendix B.4 for full details). As a second baseline, we include plain self-play RL with a minimal reward function (Appendix B.2). Figure 3 shows the results visually; Table ?? reports the numbers.

Regularized self-play RL outperforms SMART across all conditions at less than 0.25% of the human data budget. Regularized self-play outperforms SMART-tiny-CLSFT across every data bin and every metric. Plain self-play RL is already a strong baseline, achieving a human-replay at-fault collision rate of 1.4% (---). Adding as little as 30 minutes to 3 hours of human data (—●—) pulls the policy toward more human-compatible behavior and reduces the at-fault collision rate to 0.6–1%. SMART, trained on 52 days of human demonstrations, achieves an at-fault collision rate of 1.6%, which is two to three times higher than regularized self-play RL trained on less than 0.25% as much human demonstration data. Regularized self-play also achieves near-perfect task completion, while SMART agents reach their destination at most 90% of the time. Part of this gap is structural: self-play RL agents train on far more transitions than SMART, 10 billion versus 225 million (Figure 1). But the more fundamental difference is that self-play RL learns primarily from synthetic experience, not human demonstrations.

Regularized self-play agents are robust across populations. Regularized self-play agents pair well both with their own population and with the human-replay population, as reflected in low collision rates across both settings. SMART, even at 52 days of training data, incurs a 6% collision rate when all agents in the scene run the same policy, which is more than an order of magnitude higher than the 0.3-0.8% rate of regularized self-play RL agents under the same human demonstration data conditions.

4.2 The Role of Metadata in Generalization

Scenario *metadata* (road graphs, initial agent positions, and velocity) is typically cheaper than human demonstration data, and it comes with most available driving datasets [23, 60]. Recent work has shown that regularized self-play RL in simulation grounded by target-city metadata can adapt driving policies to new cities [22]. A natural follow-up question is *how much* the diversity provided by




Figure 3: **Scaling laws for human demonstration data.** Performance of *regularized* self-play RL (—●—) and SMART with CAT-K closed-loop fine-tuning (—●—) as a function of the amount of human demonstration data used for training, evaluated in self-play and against human replays. All policies are evaluated on the same 10k held-out scenarios from the Waymo Open validation set [23]. Regularized self-play saturates at roughly three hours of human data, whereas SMART continues to improve across the full WOMB training set. Self-play RL (---) is shown as a horizontal line, since it uses no human demonstrations. The horizontal axis is semi-logarithmic.

Table 1: Performance vs. amount of human demonstration data at fixed 50k self-play training maps. Top-3 values per column are highlighted (best, 2nd, 3rd); best value additionally in bold. The unregularized self-play row uses no human demonstrations. HR score for SMART is not available (—).

Human demos used	Method	Self-play (test)			Human-replay (test)				
		Coll. (%) ↓	Off-road (%) ↓	Route prog. (%) ↑	Score ↑	Coll. (%) ↓	At-fault (%) ↓	Off-road (%) ↓	Route prog. (%) ↑
10 min	SMART	11.9	55.8	84.5	0.246	32.0	25.0	18.6	57.7
30 min	SMART	9.5	55.4	85.8	0.379	17.9	12.5	16.8	76.9
3 hours	SMART	8.0	53.6	86.2	0.518	11.4	6.9	4.5	81.5
52.1 days	SMART	6.1	53.5	91.7	0.654	4.4	1.6	1.1	88.5
—	unreg. self-play	1.0 ± 0.3	0.5 ± 0.1	99.9 ± 0.0	0.965 ± 0.003	2.9 ± 0.2	1.4 ± 0.2	1.1 ± 0.1	99.9 ± 0.0
10 min	reg. self-play (ours)	0.8 ± 0.2	0.6 ± 0.1	99.0 ± 0.2	0.943 ± 0.003	4.4 ± 0.3	1.8 ± 0.2	1.6 ± 0.2	99.8 ± 0.1
30 min	reg. self-play (ours)	0.3 ± 0.1	0.5 ± 0.1	99.3 ± 0.2	0.964 ± 0.003	2.3 ± 0.2	1.0 ± 0.1	1.1 ± 0.1	99.9 ± 0.0
3 hours	reg. self-play (ours)	0.3 ± 0.1	0.4 ± 0.1	99.4 ± 0.1	0.971 ± 0.002	1.9 ± 0.2	0.6 ± 0.1	0.9 ± 0.1	99.9 ± 0.0
30 hours	reg. self-play (ours)	0.5 ± 0.2	0.4 ± 0.1	99.6 ± 0.1	0.974 ± 0.002	1.6 ± 0.2	0.7 ± 0.1	1.1 ± 0.1	100.0 ± 0.0

metadata matters for training generalizable policies, which is what we explore here. The results in the previous section assume the metadata from 50k training maps, since the focus there was on scaling laws for human demonstration data.

We train regularized and unregularized agents on nested subsets $\mathcal{M}_k$ with $|\mathcal{M}_k| \in \{10, 100, 1,000, 10,000, 50,000\}$ scenarios, holding the BC anchors τ^n and reward function r fixed. This isolates the effect of *scenario diversity*. Each policy is trained for 5 billion steps; Figure 17 confirms convergence on the training set.

To improve the signal of the comparison between policies, we evaluate on an interactive partition of the held-out set: scenarios filtered to contain dense multi-agent interactions such as merges, unprotected turns, and yielding, where coordination with other drivers is essential (details in Appendix D.1). Most scenarios in the Waymo Open validation set involve little interaction, and performance on them saturates quickly for all methods, which compresses the differences between policies. The interactive




Figure 4: **Scaling laws for scenario metadata.** The unregularized self-play baseline is shown in blue; shades of purple correspond to regularized policies trained with different BC anchors, with darker shades indicating more anchor data. **Left:** collision rate in self-play, where all agents are controlled by the same policy on a held-out validation set. **Center left:** collision rate in human-replay, where only the SDC is controlled and the remaining agents follow their logged trajectories. **Center right:** at-fault collision rate, the fraction of collisions caused by the controlled agent. This metric is more informative than raw collision rate, because replay agents frequently strike the SDC from behind, and such collisions would be unlikely in practice (See cartoon in Figure 2). **Right:** ZSC gap Δ_{ZSC} , the difference in score between the self-play and human-replay settings.

partition preserves those differences and is the more informative signal for zero-shot coordination. Figure 4 and Table ?? report results on this partition; Table 11 reports the same evaluation on the full validation set for completeness.

Scene diversity is key for generalizing to unseen maps and for zero-shot coordination with the human population. Both plain and regularized self-play improve sharply with the amount of metadata. For plain self-play, the at-fault collision rate drops from 43.8% at 10 scenarios to 2.7% at 10k scenarios, and the human-replay collision rate falls from 47.1% to 6.0% over the same range. Regularized self-play follows the same trend and reaches much lower absolute values: with a 3-hour BC anchor, the at-fault collision rate drops from 22.0% at 10 scenarios to 0.0% at 10k scenarios. The ZSC gap closes in parallel, shrinking from a self-play score that far exceeds the human-replay score to near parity at 10k scenarios (self-play 0.971, human-replay 0.945 for the 3-hour anchor). Diversity during self-play is what drives generalization to unseen maps.

Plain self-play stagnates at 2.7% at-fault collision rate. We see that metadata helps, but only up to a point. The best plain self-play policy, trained on 10k maps, reaches a human-replay collision rate of 6.0% and an at-fault collision rate of 2.7%. Scaling to 50k scenarios yields no further improvement: collision rate ticks up to 6.9% and at-fault rate to 2.8%, both within noise. The bottleneck is no longer scene diversity but the absence of a prior on human-compatible behavior. Plain self-play converges to equilibria that are internally consistent but diverge from how humans actually drive, and more map diversity cannot resolve this discrepancy.

Table 2: Scaling evaluation on held-out Waymo scenarios. Self-play metrics are reported on 10k randomly sampled validation scenarios; human-replay metrics on 200 interactive validation scenarios. Top-3 values per column are highlighted (best, 2nd, 3rd); best value additionally in bold. Gray marks the best unregularized value per column.

Self-play maps (metadata)	Anchor data (human demonstrations)	Self-play (test)			Human-replay (test)					
		Score $\uparrow$	Coll. (%) $\downarrow$	Off-road (%) $\downarrow$	Score $\uparrow$	Coll. (%) $\downarrow$	At-fault (%) $\downarrow$	Rear coll. (%) $\downarrow$	Route prog. $\uparrow$	Lateral L2 $\downarrow$
10	0 (unreg.)	0.712 $\pm$ 0.010	14.8 $\pm$ 0.8	19.3 $\pm$ 0.8	0.323 $\pm$ 0.007	48.9 $\pm$ 0.7	44.6 $\pm$ 0.7	34.3 $\pm$ 0.7	0.994 $\pm$ 0.001	4.16 $\pm$ 0.16
100	0 (unreg.)	0.898 $\pm$ 0.006	6.3 $\pm$ 0.5	5.6 $\pm$ 0.5	0.591 $\pm$ 0.007	34.3 $\pm$ 0.7	29.3 $\pm$ 0.6	19.7 $\pm$ 0.6	0.997 $\pm$ 0.001	5.53 $\pm$ 0.22
1k	0 (unreg.)	0.974 $\pm$ 0.004	1.4 $\pm$ 0.3	1.0 $\pm$ 0.2	0.877 $\pm$ 0.005	11.1 $\pm$ 0.4	7.2 $\pm$ 0.4	8.1 $\pm$ 0.4	0.998 $\pm$ 0.001	3.31 $\pm$ 0.14
10k	0 (unreg.)	0.985 $\pm$ 0.003	0.8 $\pm$ 0.2	0.4 $\pm$ 0.1	0.902 $\pm$ 0.004	9.0 $\pm$ 0.4	4.1 $\pm$ 0.3	6.8 $\pm$ 0.4	1.000 $\pm$ 0.000	1.96 $\pm$ 0.05
50k	0 (unreg.)	0.986 $\pm$ 0.003	0.7 $\pm$ 0.2	0.3 $\pm$ 0.1	0.920 $\pm$ 0.004	6.7 $\pm$ 0.3	3.8 $\pm$ 0.3	5.1 $\pm$ 0.3	0.995 $\pm$ 0.001	1.81 $\pm$ 0.06
50k	10 minutes	0.968 $\pm$ 0.004	0.7 $\pm$ 0.2	0.5 $\pm$ 0.1	0.917 $\pm$ 0.004	7.4 $\pm$ 0.4	3.2 $\pm$ 0.2	6.9 $\pm$ 0.4	0.998 $\pm$ 0.000	1.24 $\pm$ 0.03
50k	30 minutes	0.976 $\pm$ 0.004	0.5 $\pm$ 0.2	0.4 $\pm$ 0.1	0.944 $\pm$ 0.003	4.8 $\pm$ 0.3	2.5 $\pm$ 0.2	4.1 $\pm$ 0.3	1.000 $\pm$ 0.000	1.08 $\pm$ 0.02
50k	3 hours	0.984 $\pm$ 0.002	0.3 $\pm$ 0.1	0.5 $\pm$ 0.1	0.941 $\pm$ 0.003	5.3 $\pm$ 0.3	2.3 $\pm$ 0.2	4.8 $\pm$ 0.3	1.000 $\pm$ 0.000	0.93 $\pm$ 0.01
50k	30 hours	0.988 $\pm$ 0.002	0.3 $\pm$ 0.1	0.3 $\pm$ 0.1	0.934 $\pm$ 0.003	5.4 $\pm$ 0.3	2.1 $\pm$ 0.2	4.7 $\pm$ 0.3	0.999 $\pm$ 0.000	0.89 $\pm$ 0.01

Table 3: Human-replay performance on the randomly sampled validation set vs. the interactive validation subset for all scaling checkpoints. Top-3 values per column are highlighted (**best**, **2nd**, **3rd**); best value additionally in bold.

Self-play maps (metadata)	Anchor data (human demonstrations)	HR — validation (random)						HR — interactive					
		Score $\uparrow$	Coll. (%) $\downarrow$	At-fault (%) $\downarrow$	Rear coll. (%) $\downarrow$	Route prog. $\uparrow$	Lateral L2 $\downarrow$	Score $\uparrow$	Coll. (%) $\downarrow$	At-fault (%) $\downarrow$	Rear coll. (%) $\downarrow$	Route prog. $\uparrow$	Lateral L2 $\downarrow$
10	0 (unreg.)	0.559 $\pm$ 0.007	28.1 $\pm$ 0.6	25.9 $\pm$ 0.6	15.4 $\pm$ 0.5	0.994 $\pm$ 0.001	2.63 $\pm$ 0.12	0.316 $\pm$ 0.015	49.6 $\pm$ 1.6	46.3 $\pm$ 1.6	36.4 $\pm$ 1.5	0.997 $\pm$ 0.001	4.19 $\pm$ 0.36
100	0 (unreg.)	0.776 $\pm$ 0.006	18.6 $\pm$ 0.5	16.6 $\pm$ 0.5	9.6 $\pm$ 0.4	0.997 $\pm$ 0.001	3.13 $\pm$ 0.16	0.591 $\pm$ 0.015	34.8 $\pm$ 1.5	29.6 $\pm$ 1.4	20.4 $\pm$ 1.3	0.999 $\pm$ 0.001	5.41 $\pm$ 0.50
1k	0 (unreg.)	0.924 $\pm$ 0.004	5.9 $\pm$ 0.3	4.3 $\pm$ 0.3	2.6 $\pm$ 0.2	0.999 $\pm$ 0.000	1.50 $\pm$ 0.08	0.865 $\pm$ 0.011	12.4 $\pm$ 1.0	9.0 $\pm$ 0.9	8.8 $\pm$ 0.9	0.998 $\pm$ 0.001	3.11 $\pm$ 0.29
10k	0 (unreg.)	0.954 $\pm$ 0.003	3.8 $\pm$ 0.3	2.3 $\pm$ 0.2	2.0 $\pm$ 0.2	0.999 $\pm$ 0.000	1.02 $\pm$ 0.03	0.896 $\pm$ 0.010	9.4 $\pm$ 0.9	4.3 $\pm$ 0.6	6.9 $\pm$ 0.8	1.000 $\pm$ 0.000	2.04 $\pm$ 0.12
50k	0 (unreg.)	0.965 $\pm$ 0.003	2.9 $\pm$ 0.2	1.4 $\pm$ 0.2	1.9 $\pm$ 0.2	0.999 $\pm$ 0.000	1.07 $\pm$ 0.04	0.929 $\pm$ 0.008	6.5 $\pm$ 0.8	2.1 $\pm$ 0.5	5.6 $\pm$ 0.7	0.998 $\pm$ 0.001	1.67 $\pm$ 0.08
10	30 minutes	0.714 $\pm$ 0.006	16.5 $\pm$ 0.5	14.8 $\pm$ 0.5	8.0 $\pm$ 0.4	0.999 $\pm$ 0.000	0.96 $\pm$ 0.02	0.436 $\pm$ 0.016	36.4 $\pm$ 1.5	33.4 $\pm$ 1.5	20.9 $\pm$ 1.3	0.998 $\pm$ 0.001	1.75 $\pm$ 0.05
10	3 hours	0.660 $\pm$ 0.007	18.9 $\pm$ 0.5	17.5 $\pm$ 0.5	7.9 $\pm$ 0.4	0.990 $\pm$ 0.001	0.98 $\pm$ 0.03	0.370 $\pm$ 0.015	42.6 $\pm$ 1.5	38.8 $\pm$ 1.5	23.5 $\pm$ 1.3	0.982 $\pm$ 0.004	1.90 $\pm$ 0.08
100	30 minutes	0.861 $\pm$ 0.005	11.0 $\pm$ 0.4	9.1 $\pm$ 0.4	6.1 $\pm$ 0.3	1.000 $\pm$ 0.000	1.00 $\pm$ 0.03	0.710 $\pm$ 0.014	25.0 $\pm$ 1.4	19.5 $\pm$ 1.2	17.4 $\pm$ 1.2	1.000	1.85 $\pm$ 0.07
100	3 hours	0.835 $\pm$ 0.005	12.8 $\pm$ 0.5	11.3 $\pm$ 0.4	6.4 $\pm$ 0.3	1.000 $\pm$ 0.000	0.91 $\pm$ 0.03	0.665 $\pm$ 0.015	27.7 $\pm$ 1.4	23.5 $\pm$ 1.3	17.4 $\pm$ 1.2	1.000 $\pm$ 0.000	1.78 $\pm$ 0.07
1k	30 minutes	0.942 $\pm$ 0.003	4.5 $\pm$ 0.3	2.8 $\pm$ 0.2	3.1 $\pm$ 0.2	0.999 $\pm$ 0.000	0.68 $\pm$ 0.02	0.898 $\pm$ 0.009	9.0 $\pm$ 0.9	5.3 $\pm$ 0.7	7.1 $\pm$ 0.8	1.000	1.27 $\pm$ 0.04
1k	3 hours	0.949 $\pm$ 0.003	3.9 $\pm$ 0.3	2.7 $\pm$ 0.2	2.1 $\pm$ 0.2	1.000 $\pm$ 0.000	0.62 $\pm$ 0.02	0.877 $\pm$ 0.010	11.3 $\pm$ 1.0	7.5 $\pm$ 0.8	7.8 $\pm$ 0.8	0.999 $\pm$ 0.001	1.15 $\pm$ 0.05
10k	10 minutes	0.946 $\pm$ 0.003	4.1 $\pm$ 0.3	1.5 $\pm$ 0.2	3.7 $\pm$ 0.3	0.997 $\pm$ 0.001	0.75 $\pm$ 0.02	0.916 $\pm$ 0.009	6.7 $\pm$ 0.8	3.4 $\pm$ 0.6	6.0 $\pm$ 0.7	0.993 $\pm$ 0.002	1.28 $\pm$ 0.07
10k	30 minutes	0.964 $\pm$ 0.003	2.6 $\pm$ 0.2	1.3 $\pm$ 0.2	2.1 $\pm$ 0.2	0.999 $\pm$ 0.000	0.63 $\pm$ 0.01	0.925 $\pm$ 0.008	7.2 $\pm$ 0.8	2.9 $\pm$ 0.5	6.9 $\pm$ 0.8	0.998 $\pm$ 0.001	1.21 $\pm$ 0.05
10k	3 hours	0.968 $\pm$ 0.002	2.2 $\pm$ 0.2	1.0 $\pm$ 0.1	1.5 $\pm$ 0.2	0.999 $\pm$ 0.000	0.56 $\pm$ 0.01	0.899 $\pm$ 0.009	8.2 $\pm$ 0.9	3.3 $\pm$ 0.6	6.9 $\pm$ 0.8	0.997 $\pm$ 0.002	1.10 $\pm$ 0.04
10k	30 hours	0.970 $\pm$ 0.002	2.0 $\pm$ 0.2	0.9 $\pm$ 0.1	1.4 $\pm$ 0.2	1.000 $\pm$ 0.000	0.51 $\pm$ 0.01	0.930 $\pm$ 0.008	5.8 $\pm$ 0.7	3.4 $\pm$ 0.6	4.6 $\pm$ 0.7	0.998 $\pm$ 0.001	1.02 $\pm$ 0.05
50k	10 minutes	0.943 $\pm$ 0.003	4.4 $\pm$ 0.3	1.8 $\pm$ 0.2	4.0 $\pm$ 0.3	0.998 $\pm$ 0.001	0.73 $\pm$ 0.01	0.924 $\pm$ 0.008	6.6 $\pm$ 0.8	2.7 $\pm$ 0.5	6.1 $\pm$ 0.7	0.998 $\pm$ 0.001	1.21 $\pm$ 0.05
50k	30 minutes	0.964 $\pm$ 0.003	2.3 $\pm$ 0.2	1.0 $\pm$ 0.1	2.0 $\pm$ 0.2	0.999 $\pm$ 0.000	0.66 $\pm$ 0.01	0.931 $\pm$ 0.008	6.5 $\pm$ 0.8	3.5 $\pm$ 0.6	4.9 $\pm$ 0.7	0.996 $\pm$ 0.002	1.25 $\pm$ 0.05
50k	3 hours	0.971 $\pm$ 0.002	1.9 $\pm$ 0.2	0.6 $\pm$ 0.1	1.7 $\pm$ 0.2	0.999 $\pm$ 0.000	0.53 $\pm$ 0.01	0.932 $\pm$ 0.008	5.8 $\pm$ 0.7	2.1 $\pm$ 0.5	4.9 $\pm$ 0.7	1.000	0.95 $\pm$ 0.03
50k	30 hours	0.974 $\pm$ 0.002	1.6 $\pm$ 0.2	0.7 $\pm$ 0.1	1.1 $\pm$ 0.1	1.000 $\pm$ 0.000	0.48 $\pm$ 0.01	0.923 $\pm$ 0.008	6.2 $\pm$ 0.8	2.4 $\pm$ 0.5	5.4 $\pm$ 0.7	0.999 $\pm$ 0.001	0.98 $\pm$ 0.05

Table 4: Main results comparing unregularized and regularized self-play at 50k training maps. Self-play score on 10k validation scenes; human-replay and IDM-replay metrics on 200 interactive validation scenes. Best value per column in **bold**.

Method	Self-play	Score		IDM	Human-replay (interactive)			
		HR			At-fault (%) $\downarrow$	Long. L2 $\downarrow$	Lat. L2 $\downarrow$	Disp. err. $\downarrow$
Unregularized	0.984 $\pm$ 0.003	0.929 $\pm$ 0.008	0.898 $\pm$ 0.009		2.1 $\pm$ 0.5	12.614 $\pm$ 0.163	1.669 $\pm$ 0.076	12.672 $\pm$ 0.173
Regularized (ours)	0.985 $\pm$ 0.002	0.932 $\pm$ 0.008	0.896 $\pm$ 0.010		2.1 $\pm$ 0.5	6.208 $\pm$ 0.104	0.946 $\pm$ 0.030	6.424 $\pm$ 0.103

4.3 Behavior Analysis

Distributional realism via WOSAC. We analyze the Waymo Open Sim Agent Challenge (WOSAC) scores to compare the distributional realism of regularized self-play agents against vanilla self-play and SMART. Figure 5 reports the WOSAC realism meta-score alongside its three group metrics (kinematic, interactive, and map-based); Figure 6 breaks down all nine submetrics that determine the meta-score.

Vanilla self-play agents score 0.68, losing points primarily in kinematics (0.22) and the interactive category. In particular, they maintain less distance to neighboring agents compared to the human logs, and exhibit increased time-to-collision (TTC) rates, reflecting the absence of human-like spatial buffering. Regularized self-play agents reach a meta-score of 0.725, and this score is stable across the full range of human data (30 minutes to 30 hours), suggesting that a small human data budget with the existing policies is sufficient to recover these realism gains. The remaining deficit appears in kinematic metrics (0.36), specifically linear speed, and in the distance-to-nearest-object submetric, indicating that while regularization meaningfully closes the gap to human behavior, some room for improvement remains. SMART at 52 days of training data achieves the highest WOSAC score of 0.75, outperforming regularized self-play on this metric despite underperforming on collision rate and task completion across all data bins.

We should note that while WOSAC is informative, the benchmark is not perfect. WOSAC measures how closely agent behavior matches the distributional statistics of human logs, but it is not obvious that higher WOSAC scores translate directly into safer or more capable driving [61].




Figure 5: Enter Caption




Figure 6: Submetrics

Regularized agents are more social drivers. Finally, we analyze a series of videos to further study the behavioral differences between regularized and unregularized agents. While we should aim to capture most behavior difference in metrics, some things are just easier to see from videos. Here, we describe some key observations from the videos. Representative videos are shared at this anonymous page here <https://sites.google.com/view/anonymous-human-like-autonomy/>.

Regularization keeps RL policies close to human anchors. Figure 7 shows task completion and KL divergence to the anchor policy over training. Both regularized and unregularized agents converge to comparably effective strategies: they learn to navigate to their goals without colliding. Yet the action distributions underlying these strategies differ substantially. Unregularized self-play RL drifts freely through the space of effective policies and converges to a strategy far from human behavior; the KL divergence increases monotonically throughout training. Regularized self-play, by contrast, remains anchored. The agent is free to discover competent strategies, but the regularization penalty keeps those strategies close to the behavioral distribution induced by human demonstrations. This is the key idea of regularized self-play: we do not constrain *what* the agent achieves, only *how* it achieves it. The result is an agent that is both capable (solving the driving task through large-scale self-play) and compatible (staying within the distribution of strategies a human driver would recognize as natural).




Figure 7: **Regularized self-play policies stay close to the anchor distribution while unregularized self-play diverges.** Here we show that, while both unregularized (black) and regularized self-play policies converge to effective strategies (left), the action distribution underlying these strategies varies widely as indicated by the kl divergence between the observation-conditioned action distributions (right). The regularized self-play policies stay close to the anchors, whereas the unregularized policies diverge.

5 Conclusion & Limitations

We considered a series of experiments aimed at putting the mixing of human demonstration data with synthetic simulated experience on a more scientific footing. Our central finding is that a pinch of human data, roughly 30 minutes to 3 hours of demonstrations, suffices to train human-compatible driving agents, when used as a behavioral cloning anchor to regularize self-play RL. This is three orders of magnitude less than state-of-the-art imitation learning baselines, eliminates the need for reward shaping, and is achieved at a fraction of the training time. The broader implication is that, wherever simulation is cheap, human data may be best used **not** as the primary training signal but as a *lightweight anchor* that steers policies away from effective-but-alien equilibria.

This empirical result raises a deeper question that we have only touched the surface of: what is the *complementary value* of a bit of human data given a batch of synthetic experience? Can we quantify this? We can loosely intuit that self-play trains more robust agents because it covers a wider distribution of states, and indeed, our self-play agents see 5 billion transitions against 225 million for the IL baseline (Figure 1). But sheer count is a crude explanation; not all transitions are equally informative. Recent work on Epiplexity [62] takes a step toward formalizing this, conceptually defining the value of data as the amount of learnable structure it contains. Translating such notions into practical tools for selecting or filtering data remains open.

As for limitations, while regularized self-play improves on prior work in multiple dimensions, several aspects remain to be improved. First, our evaluation relies on human-replay as a proxy for deployment, and it is currently unclear to what extent performance in the self-play and human-replay settings *exactly translates* to behavior on the road; closing the sim-to-real gap remains future work. Second, our best policy reaches a near-zero at-fault collision rate, but not a zero overall collision rate: replay agents can strike the SDC from behind in situations the policy cannot reasonably avoid. What counts as a fair collision-avoidance benchmark beyond at-fault heuristics like those used in this paper is itself a difficult open question, both in industry and academia [63].

Acknowledgments

We thank the authors of CAT-K [50] for generously sharing the weights of their best SMART-tiny-CLSFT checkpoint, which we use as the imitation learning baseline throughout the paper, and their code, which we use as a baseline for the scaling law experiments. We thank Wael Doulazmi for providing the basis for the SMART evaluation code and sanity checking the baseline results. We also

thank Luke Rowe, Rodrigue de Schaetzen, and Roger Girgis for feedback on some early results and various interesting discussions on the topic of end-to-end driving and self-play. We thank Momchil Tomov for a helpful discussion on evals and metrics for evaluating human-likeness and compatibility in driving.

Daphne Cornelisse is partially supported by the Cooperative AI Foundation and a Chishiki-AI SCIFE Fellowship. [FUNDING]

References

- [1] D. Silver, A. Huang, C. J. Maddison, A. Guez, L. Sifre, G. Van Den Driessche, J. Schrittwieser, I. Antonoglou, V. Panneershelvam, M. Lanctot, et al. Mastering the game of Go with deep neural networks and tree search. *Nature*, 529(7587):484–489, 2016.
- [2] J. Perolat, B. De Vylder, D. Hennes, E. Tarassov, F. Strub, V. de Boer, P. Muller, J. T. Connor, N. Burch, T. Anthony, et al. Mastering the game of Stratego with model-free multiagent reinforcement learning. *Science*, 378(6623):990–996, 2022.
- [3] S. Sokota, E. Vinitsky, H. Hu, J. Z. Kolter, and G. Farina. Superhuman ai for stratego using self-play reinforcement learning and test-time search. *arXiv preprint arXiv:2511.07312*, 2025.
- [4] S. Kazemkhani, A. Pandya, D. Cornelisse, B. Shacklett, and E. Vinitsky. Gpudrive: Data-driven, multi-agent driving simulation at 1 million fps. *arXiv preprint arXiv:2408.01584*, 2024.
- [5] M. Cusumano-Towner, D. Hafner, A. Hertzberg, B. Huval, A. Petrenko, E. Vinitsky, E. Wijmans, T. Killian, S. Bowers, O. Sener, P. Krähenbühl, and V. Koltun. Robust autonomy emerges from self-play. *arXiv preprint arXiv:2502.03349*, 2025.
- [6] D. Cornelisse, A. Pandya, K. Joseph, J. Suárez, and E. Vinitsky. Building reliable sim driving agents by scaling self-play. *arXiv preprint arXiv:2502.14706*, 2025.
- [7] W.-J. Chang, A. Rangesh, K. Joseph, M. Strong, M. Tomizuka, Y. Hu, and W. Zhan. SPACeR: Self-play anchoring with centralized reference models. *arXiv preprint arXiv:2510.18060*, 2025.
- [8] T. Yin, Z. Mei, Z. Zheng, M. Yamane, D. Wang, J. Sceats, S. M. Bateman, L. Zha, A. Badithela, O. Shorinwa, and A. Majumdar. PlayWorld: Learning robot world models from autonomous play. *arXiv preprint arXiv:2603.09030*, 2026.
- [9] J. Z. Leibo, E. Hughes, M. Lanctot, and T. Graepel. Autocurricula and the emergence of innovation from social interaction: A manifesto for multi-agent intelligence research. *arXiv preprint arXiv:1903.00742*, 2019.
- [10] A. Bakhtin, D. Wu, A. Lerer, and N. Brown. No-press diplomacy from scratch. *Advances in Neural Information Processing Systems*, 34:18063–18074, 2021.
- [11] J. Qiu, A. Saviolo, C. Wang, M. Wang, and X. Huang. Heterogeneous self-play for realistic highway traffic simulation. 2026. URL <https://arxiv.org/abs/2604.16406>.
- [12] W. B. Knox, A. Allievi, H. Banzhaf, F. Schmitt, and P. Stone. Reward (mis) design for autonomous driving. *Artificial Intelligence*, 316:103829, 2023.
- [13] D. A. Pomerleau. Alvin: An autonomous land vehicle in a neural network. *Advances in neural information processing systems*, 1, 1988.
- [14] M. Bojarski, D. Del Testa, D. Dworakowski, B. Firner, B. Flepp, P. Goyal, L. D. Jackel, M. Monfort, U. Muller, J. Zhang, et al. End to end learning for self-driving cars. *arXiv preprint arXiv:1604.07316*, 2016.
- [15] J. Pillion, X. B. Peng, and S. Fidler. Trajenglish: Traffic modeling as next-token prediction. *arXiv preprint arXiv:2312.04535*, 2023.

- [16] M. Baniodeh, K. Goel, S. Ettinger, C. Fuertes, A. Seff, T. Shen, C. Gulino, C. Yang, G. Jerfel, D. Choe, R. Wang, V. Kallem, S. Casas, R. Al-Rfou, B. Sapp, and D. Anguelov. Scaling laws of motion forecasting and planning: A technical report. *arXiv preprint arXiv:2506.08228*, 2025.
- [17] J. Suarez. PufferLib: Making reinforcement learning libraries and environments play nice. *arXiv preprint arXiv:2406.12905*, 2024.
- [18] D. Cornelisse, S. Cheng, P. Mandavilli, J. Hunt, K. Joseph, W. Doulazmi, V. Charraut, A. Gupta, J. Suarez, and E. Vinitsky. PufferDrive: A fast and friendly driving simulator for training and evaluating RL agents, 2025. URL <https://github.com/Emerge-Lab/PufferDrive>.
- [19] H. Hu, D. J. Wu, A. Lerer, J. Foerster, and N. Brown. Human-ai coordination via human-regularized search and learning. *arXiv preprint arXiv:2210.05125*, 2022.
- [20] A. Bakhtin, D. J. Wu, A. Lerer, J. Gray, A. P. Jacob, G. Farina, A. H. Miller, and N. Brown. Mastering the game of no-press Diplomacy via human-regularized reinforcement learning and planning. In *International Conference on Learning Representations*, 2023. arXiv:2210.05492.
- [21] D. Cornelisse and E. Vinitsky. Human-compatible driving partners through data-regularized self-play reinforcement learning. In *Reinforcement Learning Journal*, 2024. arXiv:2403.19648.
- [22] Z. Wang, S. Rahmani, D. Cornelisse, B. Sarkar, A. D. Goldie, J. N. Foerster, and S. Whiteson. Learning to drive in new cities without human demonstrations. *arXiv preprint arXiv:2602.15891*, 2026.
- [23] S. Ettinger, S. Cheng, B. Caine, C. Liu, H. Zhao, S. Pradhan, Y. Chai, B. Sapp, C. R. Qi, Y. Zhou, et al. Large scale interactive motion forecasting for autonomous driving: The waymo open motion dataset. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 9710–9719, 2021.
- [24] L. Chen, P. Wu, K. Chitta, B. Jaeger, A. Geiger, and H. Li. End-to-end autonomous driving: Challenges and frontiers. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 46(12):10164–10183, 2024.
- [25] X. Jia, Z. Yang, Q. Li, Z. Zhang, and J. Yan. Bench2drive: Towards multi-ability benchmarking of closed-loop end-to-end autonomous driving. *Advances in Neural Information Processing Systems*, 37:819–844, 2024.
- [26] Y. Hu, J. Yang, L. Chen, K. Li, C. Sima, X. Zhu, S. Chai, S. Du, T. Lin, W. Wang, et al. Planning-oriented autonomous driving. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 17853–17862, 2023.
- [27] B. Jiang, S. Chen, Q. Xu, B. Liao, J. Chen, H. Zhou, Q. Zhang, W. Liu, C. Huang, and X. Wang. Vad: Vectorized scene representation for efficient autonomous driving. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 8340–8350, 2023.
- [28] Y. Huang, J. Du, Z. Yang, Z. Zhou, L. Zhang, and H. Chen. A survey on trajectory-prediction methods for autonomous driving. *IEEE transactions on intelligent vehicles*, 7(3):652–674, 2022.
- [29] N. Montali, J. Lambert, P. Mougin, A. Kuefler, N. Rhinehart, M. Li, C. Gulino, T. Emrich, Z. Yang, S. Whiteson, et al. The waymo open sim agents challenge. *Advances in Neural Information Processing Systems*, 36:59151–59171, 2023.
- [30] E. Vinitsky, N. Lichtlé, S. Kanaa, et al. Nocturne: a scalable driving benchmark for bringing multi-agent learning one step closer to the real world. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- [31] H. Caesar, V. Bankiti, A. Lang, S. Vora, V. Liong, Q. Xu, A. Krishnan, Y. Pan, G. Baldan, and O. Beijbom. nuscenes: A multimodal dataset for autonomous driving. arxiv. 2019.

- [32] B. Wilson, W. Qi, T. Agarwal, J. Lambert, J. Singh, S. Khandelwal, B. Pan, R. Kumar, A. Hartnett, J. K. Pontes, et al. Argoverse 2: Next generation datasets for self-driving perception and forecasting. *arXiv preprint arXiv:2301.00493*, 2023.
- [33] M. Bansal, A. Krizhevsky, and A. Ogale. Chauffeurnet: Learning to drive by imitating the best and synthesizing the worst. *arXiv preprint arXiv:1812.03079*, 2018.
- [34] T. Salzmann, B. Ivanovic, P. Chakravarty, and M. Pavone. Trajectron++: Dynamically-feasible trajectory forecasting with heterogeneous data. In *Computer Vision–ECCV 2020: 16th European Conference, Glasgow, UK, August 23–28, 2020, Proceedings, Part XVIII 16*, pages 683–700. Springer, 2020.
- [35] J. Gu, C. Sun, and H. Zhao. Densetnt: End-to-end trajectory prediction from dense goal sets. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 15303–15312, 2021.
- [36] N. Nayakanti, R. Al-Rfou, A. Zhou, K. Goel, K. S. Refaat, and B. Sapp. Wayformer: Motion forecasting via simple & efficient attention networks. In *2023 IEEE International Conference on Robotics and Automation (ICRA)*, pages 2980–2987. IEEE, 2023.
- [37] J. Ngiam, B. Caine, V. Vasudevan, Z. Zhang, H.-T. L. Chiang, J. Ling, R. Roelofs, A. Bewley, C. Liu, A. Venugopal, et al. Scene transformer: A unified architecture for predicting multiple agent trajectories. *arXiv preprint arXiv:2106.08417*, 2021.
- [38] Z. Zhou, L. Ye, J. Wang, K. Wu, and K. Lu. Hivt: Hierarchical vector transformer for multi-agent motion prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 8823–8833, 2022.
- [39] S. Shi, L. Jiang, D. Dai, and B. Schiele. Motion transformer with global intention localization and local movement refinement. *Advances in Neural Information Processing Systems*, 35: 6531–6543, 2022.
- [40] Z. Zhou, J. Wang, Y.-H. Li, and Y.-K. Huang. Query-centric trajectory prediction. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 17863–17873, 2023.
- [41] A. Seff, B. Cera, D. Chen, M. Ng, A. Zhou, N. Nayakanti, K. S. Refaat, R. Al-Rfou, and B. Sapp. Motionlm: Multi-agent motion forecasting as language modeling. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pages 8579–8590, 2023.
- [42] W. Wu, X. Feng, Z. Gao, and Y. Kan. Smart: Scalable multi-agent real-time motion generation via next-token prediction. *Advances in Neural Information Processing Systems*, 37:114048–114071, 2024.
- [43] Z. Zhong, D. Rempe, D. Xu, Y. Chen, S. Veer, T. Che, B. Ray, and M. Pavone. Guided conditional diffusion for controllable traffic simulation. In *2023 IEEE international conference on robotics and automation (ICRA)*, pages 3560–3566. IEEE, 2023.
- [44] C. M. Jiang, Y. Bai, A. Cornman, C. Davis, X. Huang, H. Jeon, S. Kulshrestha, J. Lambert, S. Li, X. Zhou, et al. Scenediffuser: Efficient and controllable driving simulation initialization and rollout. *Advances in Neural Information Processing Systems*, 37:55729–55760, 2024.
- [45] Z. Huang, Z. Zhang, A. Vaidya, Y. Chen, C. Lv, and J. F. Fisac. Versatile behavior diffusion for generalized traffic agent simulation. *arXiv preprint arXiv:2404.02524*, 2024.
- [46] S. Tan, J. Lambert, H. Jeon, S. Kulshrestha, Y. Bai, J. Luo, D. Anguelov, M. Tan, and C. M. Jiang. Scenediffuser++: City-scale traffic simulation via a generative world model. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 1570–1580, 2025.

- [47] B. Liao, S. Chen, H. Yin, B. Jiang, C. Wang, S. Yan, X. Zhang, X. Li, Y. Zhang, Q. Zhang, et al. Diffusiondrive: Truncated diffusion model for end-to-end autonomous driving. In *Proceedings of the Computer Vision and Pattern Recognition Conference*, pages 12037–12047, 2025.
- [48] Y. Lu, J. Fu, G. Tucker, X. Pan, E. Bronstein, R. Roelofs, B. Sapp, B. White, A. Faust, S. Whiteson, et al. Imitation is not enough: Robustifying imitation with reinforcement learning for challenging driving scenarios. In *2023 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 7553–7560. IEEE, 2023.
- [49] Z. Peng, W. Luo, Y. Lu, T. Shen, C. Gulino, A. Seff, and J. Fu. Improving agent behaviors with RL fine-tuning for autonomous driving. In *Computer Vision - ECCV 2024 - 18th European Conference*, volume 15083 of *Lecture Notes in Computer Science*, pages 165–181. Springer, 2024.
- [50] Z. Zhang, P. Karkus, M. Igl, W. Ding, Y. Chen, B. Ivanovic, and M. Pavone. Closed-loop supervised fine-tuning of tokenized traffic models. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2025.
- [51] D. Silver, T. Hubert, J. Schrittwieser, I. Antonoglou, M. Lai, A. Guez, M. Lanctot, L. Sifre, D. Kumaran, T. Graepel, et al. A general reinforcement learning algorithm that masters chess, shogi, and go through self-play. *Science*, 362(6419):1140–1144, 2018.
- [52] O. Vinyals, I. Babuschkin, W. M. Czarnecki, M. Mathieu, A. Dudzik, J. Chung, D. H. Choi, R. Powell, T. Ewalds, P. Georgiev, et al. Grandmaster level in starcraft ii using multi-agent reinforcement learning. *nature*, 575(7782):350–354, 2019.
- [53] H. Hu, A. Lerer, A. Peysakhovich, and J. Foerster. “other-play” for zero-shot coordination. In *International conference on machine learning*, pages 4399–4410. PMLR, 2020.
- [54] N. Bard, J. N. Foerster, S. Chandar, N. Burch, M. Lanctot, H. F. Song, E. Parisotto, V. Dumoulin, S. Moitra, E. Hughes, et al. The hanabi challenge: A new frontier for ai research. *Artificial Intelligence*, 280:103216, 2020.
- [55] A. P. Jacob, D. J. Wu, G. Farina, A. Lerer, H. Hu, A. Bakhtin, J. Andreas, and N. Brown. Modeling strong and human-like gameplay with KL-regularized search. In *International Conference on Machine Learning*, pages 9695–9728. PMLR, 2022.
- [56] M. Treiber, A. Hennecke, and D. Helbing. Congested traffic states in empirical observations and microscopic simulations. *Physical Review E*, 62(2):1805–1824, Aug. 2000. ISSN 1063-651X, 1095-3787. doi:10.1103/PhysRevE.62.1805. URL <https://link.aps.org/doi/10.1103/PhysRevE.62.1805>.
- [57] V. Charraut, W. Doulazmi, T. Tournaire, and T. Buhet. V-Max: A RL framework for autonomous driving. *Reinforcement Learning Journal*, 6:2427–2451, 2025.
- [58] D. Dauner, M. Hallgarten, T. Li, X. Weng, Z. Huang, Z. Yang, H. Li, I. Gilitschenski, B. Ivanovic, M. Pavone, et al. Navsim: Data-driven non-reactive autonomous vehicle simulation and benchmarking. *Advances in Neural Information Processing Systems*, 37:28706–28719, 2024.
- [59] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov. Proximal policy optimization algorithms. *arXiv preprint arXiv:1707.06347*, 2017.
- [60] H. Caesar, J. Kabzan, K. S. Tan, W. K. Fong, E. Wolff, A. Lang, L. Fletcher, O. Beijbom, and S. Omari. nuplan: A closed-loop ml-based planning benchmark for autonomous vehicles. *arXiv preprint arXiv:2106.11810*, 2021.
- [61] D. Cornelisse. Human-likeness metrics for autonomous agents: are we measuring the right thing? Substack, 2025. Blog post analyzing the Waymo Open Sim Agent Challenge (WOSAC) realism benchmark.

- [62] M. Finzi, S. Qiu, Y. Jiang, P. Izmailov, J. Z. Kolter, and A. G. Wilson. From entropy to epiplexity: Rethinking information for computationally bounded intelligence. *arXiv preprint arXiv:2601.03220*, 2026.
- [63] J. M. Scanlon, K. D. Kusano, J. Engstrom, and T. Victor. Collision avoidance effectiveness of an automated driving system using a human driver behavior reference model in reconstructed fatal collisions. In *WCX SAE World Congress Experience*. SAE Technical Paper, 2026.
- [64] Waymo LLC. Waymo safety impact. <https://waymo.com/safety/impact/>, 2025. Accessed: 2026-05-06.
- [65] C. Gulino, J. Fu, W. Luo, G. Tucker, E. Bronstein, Y. Lu, J. Harb, X. Pan, Y. Wang, X. Chen, et al. Waymax: An accelerated, data-driven simulator for large-scale autonomous driving research. *Advances in Neural Information Processing Systems*, 36:7730–7742, 2023.

A Simulation Environment and Design

A.1 World Initialization from Scenario Metadata

We use PufferDrive 2.0 for simulation and training [18]. PufferDrive is a batched simulator that runs many environments in parallel, reaching 400k steps per second (SPS) on an NVIDIA RTX 5090 GPU. We initialize environments using the Waymo Open Motion Dataset (WOMD) [23], which provides a large set of multi-agent traffic scenarios. Each scenario supplies the metadata we need: the roadgraph, a variable number of agents (cars, cyclists, and pedestrians), and other objects in the scene. This information is the output of a perception stack, so we operate directly on these clean features (in bounding-box world).

Each scenario is 9 seconds long and discretized into 90 steps. We take each logged agent’s initial position ($t = 0$) as its starting position in the scene, and its last valid logged position ($t = T$) as its goal, which lets us goal-condition the agents. The full Waymo training dataset contains 500k scenarios, but in this paper we use at most 50k of the randomly sampled scenarios. When constructing the environments, we randomly sample scenarios from WOMD until we hit a target number of agents (e.g., on an NVIDIA RTX 4080 with 16GB of memory, we keep adding environments until we reach 1024 agents).

A.2 Observation Space

We take a decentralized approach and provide every agent with a partial view of the environment in a local coordinate frame. This is similar to the observation space of prior related works, such as GIGAFLOW [5], and GPUDrive [4]. At each timestep, an agent receives the combination of three feature blocks: an ego block describing its own state, a partner block describing the $N_p = 31$ closest other agents within a 50 m radius, and a road block describing up to $N_r = 128$ nearby road segments drawn from a 21×21 grid of $5 \text{ m} \times 5 \text{ m}$ cells centered on the agent. Missing slots (fewer partners or road segments than the maximum) are zero-padded. Tables 5, 6, and 7 list the features in each block. All positions and headings are expressed in the agent’s local frame, so the observation is invariant to the global pose of the scene. The total observation vector has dimension $14 + 7 \times 31 + 7 \times 128 = 1,127$.

Table 5: Ego features (14 values) for the delta-local dynamics model. Features 0–3 expose the sampled conditioning variables to the policy so it can modulate its behavior as a function of λ and the reward weights (Section 4.2). Features 11–13 expose the previously executed (post-constraint) action, which gives the policy access to the kinematic state needed to respect the acceleration and yaw-rate bounds described in Section A.3.

Idx	Feature	Normalization	Description
0	λ	—	Human-regularization coefficient
1	r_{coll}	—	Sampled collision reward
2	r_{off}	—	Sampled off-road reward
3	r_{goal}	—	Sampled goal reward
4	Δx_{goal}	$\times 0.005$	Goal position (ego frame), longitudinal
5	Δy_{goal}	$\times 0.005$	Goal position (ego frame), lateral
6	signed speed	$/ 100 \text{ m/s}$	Speed projected onto heading
7	vehicle width	$/ 15 \text{ m}$	Ego bounding-box width
8	vehicle length	$/ 30 \text{ m}$	Ego bounding-box length
9	collision flag	$\{0, 1\}$	1 if currently colliding
10	entity type	$/ 3$	Vehicle (1), pedestrian (2), cyclist (3)
11	prev. action Δx	$/ 3.5 \text{ m}$	Previously executed longitudinal step
12	prev. action Δy	$/ 0.1 \text{ m}$	Previously executed lateral step
13	prev. action $\Delta \psi$	$/ (\pi/6)$	Previously executed heading change

Table 6: Partner features (7 values $\times$ 31 partners = 217 values). Partners are ordered by index and filtered to those within 50 m of the ego agent. All positions and headings are in the ego frame.

Idx	Feature	Normalization	Description
0	Δx	$\times 0.02$	Partner position, longitudinal
1	Δy	$\times 0.02$	Partner position, lateral
2	partner width	/ 15 m	Partner bounding-box width
3	partner length	/ 30 m	Partner bounding-box length
4	$\cos(\Delta\psi)$	—	Relative heading, cosine component
5	$\sin(\Delta\psi)$	—	Relative heading, sine component
6	partner signed speed	/ 100 m/s	Signed speed along partner’s heading

Table 7: Road-segment features (7 values $\times$ 128 segments = 896 values). Segments are drawn from a 21×21 grid of 5 m cells centered on the ego agent, and include road lanes, road lines, and road edges. Each segment is described by the midpoint, length, and orientation of a single polyline segment.

Idx	Feature	Normalization	Description
0	midpoint x	$\times 0.02$	Segment midpoint, longitudinal (ego frame)
1	midpoint y	$\times 0.02$	Segment midpoint, lateral (ego frame)
2	segment length	/ 100 m	Length of the polyline segment
3	segment width	/ 100 m	Fixed nominal width (0.1 m)
4	$\cos(\theta)$	—	Segment orientation in ego frame
5	$\sin(\theta)$	—	Segment orientation in ego frame
6	segment type	$\{0, 1, 2\}$	Road lane (0), road line (1), road edge (2)

A.3 Actions and Dynamics

We use a single dynamics model with a discretized action space.

Delta-local dynamics with kinematic constraints. The action is a triple $(\Delta x, \Delta y, \Delta\psi)$ in the agent’s local frame at time t . Translation is rotated into the world frame and added to the position; heading is updated directly:

$$x_{t+1} = x_t + \cos(\psi_t) \Delta x - \sin(\psi_t) \Delta y, \quad (3)$$

$$y_{t+1} = y_t + \sin(\psi_t) \Delta x + \cos(\psi_t) \Delta y, \quad (4)$$

$$\psi_{t+1} = \text{wrap}(\psi_t + \Delta\psi). \quad (5)$$

Velocity is reported as the world-frame displacement divided by $\Delta t = 0.1$ s. We bound each component roughly based on realistic actions present in the human demonstrations, as shown in Figure 8; specifically, we define $\Delta x \in [-3.5, 3.5]$ m, $\Delta y \in [-0.1, 0.1]$ m, and $\Delta\psi \in [-\pi/6, \pi/6]$. Each of the three dimensions is binned independently into 512 values, yielding step sizes of approximately 0.0137 m, 3.91×10^{-4} m, and 2.05×10^{-3} rad, respectively. The action distributions for Δy and $\Delta\psi$ are roughly symmetric, whereas the distribution for Δx is strongly asymmetric. This is expected, since most vehicles move forward and only a small number of agents in the scenes drive in reverse (e.g., when parking).

Delta-local dynamics are kinematically unconstrained by default: the agent can translate laterally without rotating, pivot in place, or instantaneously reverse its heading rate. To prevent these unrealistic behaviors, we apply four physics-based constraints to the action at each step, in sequence. Each constraint clips the action after the previous one has been applied, with the previously executed (post-constraint) values used as the reference. The constraints are:

1. **Longitudinal acceleration bound.** The change in implied forward speed is clipped to $\pm A_{\text{long,max}} \cdot \Delta t$, where $A_{\text{long,max}} = 8 \text{ m/s}^2$. This caps acceleration and braking.




Figure 8: Discretized delta-local action space for each component (Δx , Δy , $\Delta \psi$). Histograms show the empirical density (blue) of 10,996,751 valid action timesteps recovered from expert trajectories across 10,000 maps. Yellow lines mark the 1st and 99th percentiles of the data; red lines mark the action-space bounds (± 3.5 m, ± 0.1 m, $\pm \pi/6$ rad). Each dimension is binned independently into 512 values. The bounds were chosen to respect natural movements in the data: 0.00% of Δx and Δy samples fall outside their bounds, and 0.71% of $\Delta \psi$ samples fall outside $\pm \pi/6$.

2. **Lateral motion envelope.** Lateral displacement is bounded by $|\Delta y| \leq |\Delta x| \cdot \tan(\delta_{\max})$, where $\delta_{\max} = 0.7$ rad is the maximum effective steering angle. This eliminates lateral sliding and side-shimmy at low forward speed.

These physical constraints prevent kinematically implausible actions; they do not encode any preference over driving style and are independent of the human anchor.

A.4 Collecting Human Demonstrations

The behavioral cloning (BC) anchor is trained on observation–action pairs (o_t, a_t) . We therefore need actions that (i) live in the simulator’s action space and (ii) reproduce the logged motion when applied through the simulator’s dynamics. We construct the dataset in two steps. Figure ?? shows three examples of this process in the simulator.

Step 1: Inferring actions from the logs. For each timestep t , we invert the delta-local dynamics to recover the action that produced the next logged state. Projecting the world-frame displacement into the agent’s local frame at t gives:

$$\Delta x_t = \cos(\psi_t)(x_{t+1} - x_t) + \sin(\psi_t)(y_{t+1} - y_t), \quad (6)$$

$$\Delta y_t = -\sin(\psi_t)(x_{t+1} - x_t) + \cos(\psi_t)(y_{t+1} - y_t), \quad (7)$$

$$\Delta \psi_t = \text{wrap}(\psi_{t+1} - \psi_t). \quad (8)$$

Each triple is clipped to the action bounds and snapped to the nearest discrete bin. Timesteps where either t or $t + 1$ is flagged invalid in the log are marked as missing and excluded from training.

Step 2: Replaying actions through the simulator. To produce observations, we replay the inferred action sequence through the simulator and record the observation it computes at every simulator state. The BC anchor is then trained on the resulting (simulator observation, inferred action) tuples.

Effect of discretization on performance. Figure 9 and Table 8 quantify the cost of discretization. Continuous actions reproduce the logged trajectory almost exactly (ADE 0.001 m), confirming that the delta-local dynamics and kinematic constraints are themselves well-posed. Discretizing into 512 bins per dimension introduces a quantization floor of ADE 0.097 m, which is roughly two orders of magnitude larger, but is still very close to the original trajectory. Off-road and collision rates increase modestly under discretization (1.2% vs. 0.8% off-road, 0.4% vs. 0.0% collision), reflecting the rare cases where snapping to the nearest bin pushes the SDC just outside a road edge or into a static neighbor; both representations complete the route in 100% of scenarios.

Table 8: Inferred-expert-action quality for the delta-local dynamics model. Comparison of discrete (bin-quantized) vs continuous (direct float) expert actions. Aggregated over 10,240 pooled samples. Values are mean $\pm$ SE.

Action type	Route prog. (%) $\uparrow$	Coll. (%) $\downarrow$	Off-road (%) $\downarrow$	ADE (m) $\downarrow$	Lateral L2 (m) $\downarrow$	Longitudinal L2 (m) $\downarrow$
discrete	100.0	0.4 ± 0.1	1.2 ± 0.2	0.097 ± 0.002	0.096 ± 0.002	0.004 ± 0.000
continuous	100.0	0.0	0.8 ± 0.1	0.001 ± 0.000	0.001 ± 0.000	0.001 ± 0.000




Figure 9: Effect of action discretization on inferred-expert-action quality. We replay each agent’s logged trajectory through the simulator using actions inferred from the logs, comparing discrete (bin-quantized, blue) and continuous (direct float, green) action representations. **Left:** SDC rates aggregated across 10,240 pooled samples; both representations complete the route in 100% of scenarios, but discretization induces modestly higher off-road and collision rates. **Center, right:** distributions of per-trajectory lateral and longitudinal L2 error to the logged pose. Continuous actions reproduce the log almost exactly (errors concentrated near zero), while discrete actions exhibit a small but consistent quantization floor of ~ 0.1 m laterally. Error bars on the bar plot denote standard error.




Figure 10: Three annotated example scenarios illustrating the human demonstration data collection process. The self-driving car (SDC), marked in cyan, is the Waymo vehicle whose human-driven trajectory we use as the demonstration. Logged trajectories are shown in green; purple trajectories show the result of stepping each agent through the simulator under the inferred delta-local actions. We select only the SDC trajectory because it is typically the cleanest demonstration in the scene; the visualized step-wise displacement illustrates a few low-quality (high-ADE) log trajectories that would otherwise contaminate the anchor.

A.5 Reward Function

We use a sparse reward: $r^i = +1$ if agent i reaches its goal within $\delta = 2$ meters before the episode ends, -1 on collision or going off-road, and 0 otherwise. We deliberately omit dense shaping terms so that safe and human-compatible behaviors can emerge from regularization.

B Training

B.1 Behavioral Cloning Anchor Policies

Each anchor τ_n is trained by minimizing the negative log-likelihood of the logged actions under the factorized discrete action distribution described in Appendix A.3. We extract observation, action tuples through the procedure described in Appendix A.4. Note that we use only the SDC trajectory from each scenario for training, as it is the highest-quality demonstration. Since other agents are reconstructed from the perception stack, they exhibit more noise. Moreover, we have no guarantees about the driving quality of the surrounding humans. Since we obtain one trajectory per scene, each scenario contributes roughly 9 seconds of human driving data. Although these trajectories were collected in Waymo vehicles, they reflect manual human driving by an expert driver behind the wheel [23].

We train with Adam at a learning rate of 10^{-4} and a batch size of 2048 for up to 5000 epochs, with early stopping on the held-out validation loss after 100 epochs without improvement. Table 9 reports open- and closed-loop metrics for each anchor on 10,000 held-out validation scenarios. Figure 12 shows the 5-bin validation accuracy for each action head over training; from only 30 minutes of data, validation accuracy converges to between 80% and 90%. This metric is informative since there are 256 bins per action head, so the step sizes are very small.

Figures 13 and 14 compare the learned action distributions against the empirical distribution of the logged actions, for anchors trained on 10 minutes (67 maps) and 3 hours (1200 maps) of data, respectively; in both cases, the learned distributions match the data reasonably well.

Table 9: BC anchor evaluation. Open-loop metrics on the held-out validation set; closed-loop metrics averaged over validation scenes. Within-5-bin accuracy is the average of Δx , Δy , Δyaw accuracies at the final training step.

Human data (h)	Acc. (%)	Open-loop		Loss	Closed-loop self-play		Closed-loop human-replay (SDC only)	
		Acc. (%)	Acc. ± 5 bins (%)		Route prog.	Score	Route prog.	Score
0.2	23.4		72.4	15.677	0.720 ± 0.012	0.215 ± 0.013	0.765 ± 0.007	0.242 ± 0.009
0.5	36.1		87.3	5.269	0.719 ± 0.011	0.277 ± 0.014	0.800 ± 0.006	0.371 ± 0.011
3.0	48.2		92.6	1.641	0.835 ± 0.010	0.502 ± 0.017	0.842 ± 0.006	0.538 ± 0.011
30.0	52.8		94.9	1.266	0.932 ± 0.007	0.685 ± 0.016	0.873 ± 0.006	0.666 ± 0.010




Figure 11: Open- and closed-loop performance of the anchor BC policies as a function of human demonstration data. Left: The final real (blue) and within 5 bin accuracy (purple) accuracy on 10,000 held-out validation scenarios. Right: Final validation loss. Right; Route progress; Right Score.

BC Training Data vs Learned Policy — 23296 samples, delta_local, 67 maps




Figure 13: Example of actual vs. learned distributions - for 67 maps (10 min)




Figure 12: Training curves for the anchor policies. Each panel shows within-5-bin validation accuracy on a held-out set of scenarios for one action component (Δx , Δy , $\Delta \psi$). Curves terminate at different step counts because training stops once validation accuracy plateaus (no improvement for 100 consecutive epochs).




Figure 14: Example of actual vs. learned distributions - for 1200 maps (3h)

B.2 Self-Play Reinforcement Learning

B.3 λ -Conditioned Regularized Self-Play RL

Our approach combines two ideas from prior work. Hu et al. [19] show that regularizing an RL policy toward a BC anchor and conditioning the policy on the regularization strength allows a single policy to express a range of human-compatible skill levels. Cusumano-Towner et al. [5] show that conditioning a policy on its sampled reward weights, rather than on a summary statistic, is an effective way to distill a diverse population of behaviors into a single steerable policy. We combine these ideas: we sample a regularization weight λ per episode from a uniform distribution and expose its value to the policy as part of the observation. This induces a heterogeneous population during self-play: agents must coordinate with partners of varying human-likeness, while keeping the trained policy steerable at inference without retraining.

B.3.1 Objective

Let π_θ denote the RL policy with parameters θ , and let τ_n be the fixed BC anchor trained on n scenarios.

Sampling. At the start of each episode, every agent i is assigned a regularization weight

$$\lambda_i \sim \text{Uniform}(0, 1), \quad (9)$$

held fixed for the duration of the episode and resampled on reset.

Regularization loss. We regularize π_θ toward τ_n by penalizing the KL divergence between the two policies on states visited during the rollout. For a minibatch of M transitions $\{(o_j, \lambda_j)\}_{j=1}^M$:

$$\mathcal{L}_{\text{reg}}(\theta) = \frac{1}{M} \sum_{j=1}^M \lambda_j D_{\text{KL}}(\tau_n(\cdot | o_j) \parallel \pi_\theta(\cdot | o_j, \lambda_j)), \quad (10)$$

where

$$D_{\text{KL}}(\tau_n \parallel \pi_\theta) = \sum_a \tau_n(a | o) [\log \tau_n(a | o) - \log \pi_\theta(a | o, \lambda)]. \quad (11)$$

Transitions with larger λ_j are pulled more strongly toward the anchor; when $\lambda_j = 0$ the regularization vanishes and the policy learns from the PPO objective alone. Because λ_j is part of the observation, π_θ learns to modulate its behavior as a continuous function of λ .

Full training objective. The standard PPO objective combines a clipped surrogate policy-gradient loss $\mathcal{L}_{\text{pg}}$, a value-function loss $\mathcal{L}_v$, and an entropy bonus H . The full training objective is

$$\mathcal{L}(\theta) = \mathcal{L}_{\text{pg}} + c_v \mathcal{L}_v - c_H H + \mathcal{L}_{\text{reg}}, \quad (12)$$

where c_v and c_H are the value and entropy coefficients. No additional scalar multiplies $\mathcal{L}_{\text{reg}}$: the per-sample weighting by λ_j already controls the regularization strength.

Inference. At test time, λ can be set to any value in $[0, 1]$ to continuously control the degree of human-likeness: $\lambda = 1$ produces the most human-like behavior, and $\lambda = 0$ recovers the unregularized self-play policy.

B.3.2 Effect of Regularization on the Policy Gradient

Consider the gradient of the regularization loss with respect to θ .

Per-sample gradient. For a single rollout transition with observation o and conditioning value λ , the regularization loss contributes

$$\nabla_\theta \lambda D_{\text{KL}}(\tau_n(\cdot | o) \parallel \pi_\theta(\cdot | o, \lambda)) = -\lambda \sum_a \tau_n(a | o) \nabla_\theta \log \pi_\theta(a | o, \lambda). \quad (13)$$

This is a weighted behavioral cloning gradient: each action a is reinforced in proportion to the anchor’s probability $\tau_n(a | o)$, scaled by λ . When λ is large, the gradient strongly pushes π_θ toward the anchor; when $\lambda = 0$, the regularization term vanishes and only the PPO objective drives learning.

Interaction with the PPO gradients. The PPO surrogate loss produces a gradient of the form

$$\nabla_\theta \mathcal{L}_{\text{pg}} \propto -\hat{A}(o, a) \nabla_\theta \log \pi_\theta(a | o, \lambda), \quad (14)$$

where $\hat{A}$ is the estimated advantage and a is the action taken during the rollout. The entropy bonus contributes

$$-c_H \nabla_\theta H[\pi_\theta] = c_H \sum_a (1 + \log \pi_\theta(a | o, \lambda)) \nabla_\theta \log \pi_\theta(a | o, \lambda). \quad (15)$$

Combining all three terms gives

$$\underbrace{-\hat{A}(o, a) \nabla_{\theta} \log \pi_{\theta}(a \mid o, \lambda)}_{\text{policy gradient}} \quad (16)$$

$$+ \underbrace{c_H \sum_{a'} (1 + \log \pi_{\theta}(a' \mid o, \lambda)) \nabla_{\theta} \log \pi_{\theta}(a' \mid o, \lambda)}_{\text{entropy bonus}} \quad (17)$$

$$- \underbrace{\lambda \sum_{a'} \tau_n(a' \mid o) \nabla_{\theta} \log \pi_{\theta}(a' \mid o, \lambda)}_{\text{KL regularization}}. \quad (18)$$

Entropy bonus vs. KL regularization. The entropy bonus and the KL term are structurally similar: both sum over the full action space and both act independently of the collected rewards. However, they pull the policy in different directions. The entropy bonus weights each action by $1 + \log \pi_{\theta}(a)$, pushing probability mass toward the uniform distribution — an *untargeted* form of exploration. The KL term weights each action by $\tau_n(a)$, pushing mass toward the anchor’s preferred actions — a *targeted* form of regularization that steers the policy toward a reference distribution learned from human data.

B.3.3 Baseline

We train the unregularized self-play baselines ($\lambda = 0$ throughout training) for 5 billion steps each on $\{10, 100, 1,000, 10,000, 50,000\}$ maps, which is sufficient for convergence (Figure 17).

B.3.4 Data Filtering

We train only on transitions from agents that have not yet reached their goal. Once an agent is within δ of its goal, it is removed from the scene and its remaining transitions are excluded from training.

B.3.5 Hyperparameters

We use PufferLib [17] for training. Hyperparameters are listed in Table 10. We use the native Protein software for sweeps.

Table 10: PPO training hyperparameters.

Architecture		Training		Environment & Rewards	
Input size	64	Total timesteps	10B	Number of agents	1,024
Hidden size	256	Batch size	524,288	Number of workers	16
RNN type	LSTM	Minibatch size	32,768	Episode length	91 steps
RNN input size	256	Rollout horizon	32	Timestep Δt	0.1 s
RNN hidden size	256	Update epochs	1	Goal radius	2.5 m
		Learning rate	4.26×10^{-3}	Action space	Discrete
		LR schedule	Linear annealing	Dynamics model	Delta-local
		Adam β_1	0.9	Goal reward	+1.0
		Adam β_2	0.999	Collision penalty	−1.0
		Adam ϵ	10^{-8}	Off-road penalty	−1.0
		Clip coefficient	0.2		
		Entropy coefficient	0.001		
		VF coefficient	2.0		
		VF clip	0.2		
		GAE λ	0.95		
		Discount γ	0.99		
		Max gradient norm	1.0		
		Priority α	0.85		
		Priority β_0	0.85		
		V-trace c clip	1.0		
		V-trace ρ clip	1.0		
		Optimizer	Adam		
		Seed	42		

B.4 SMART Model Training

We trained SMART models in CATK on subsets of the Waymo Open Motion Dataset (WOMD). WOMD motion shards were preprocessed into CATK’s per-scenario cached format, and all training subsets were built from these cached scenario files. Our final local runs used `smart_mini_3M` with vehicle-only supervision on deterministic subsets of 67, 200, 1200, and 12000 scenarios. In the subset construction scripts we used, scenarios are sorted by cached scenario filename in lexicographic order. Vehicle-only supervision means that only vehicle agents contributed to the training loss, while pedestrians and cyclists were still present in the scene and available as context to the model. The local models were trained with CATK’s `pre_bc` configuration on a single GPU for 64 epochs with batch size 8. These runs used standard supervised behavioral cloning and were not further fine-tuned with CAT-K / CLSFT. We additionally compare against two author-provided checkpoints: a behavioral-cloning checkpoint (`pre_bc_E31.ckpt`) and a closed-loop supervised fine-tuned checkpoint (`clsft_E9.ckpt`). For downstream evaluation, we exported predictions as `.pk1` files on the 10k validation scenes from the `pufferdrive_womd_subsets` dataset. We used two export modes: an all-agents mode and a planning mode in which only the SDC is controlled by the model and all other agents are replayed from ground truth. We re-exported all model and export-mode combinations with 32 rollouts for multimodal evaluation. We verified that the scenario-ID overlap between each local training subset (67, 200, 1200, and 12000) and the evaluation set is zero.

C Neural Network Architecture

Both the BC anchor and the RL policy share the same multi-modal encoder structure. The flattened environment observation vector is first unpacked into its modalities: ego state, partner agents, and road segments. Each modality is processed by a two-layer MLP with ReLU activation and layer normalisation between the two linear layers. Partner and road embeddings are then aggregated across objects via max-pooling, producing one vector per stream. The three pooled vectors are concatenated and passed through a shared two-layer MLP (Linear $\rightarrow$ ReLU $\rightarrow$ Linear) to produce the final embedding. Separate linear heads decode this embedding into logits over each action dimension; a separate linear head with unit output produces the value estimate. The two architectures differ in width and in the presence of recurrence:

- **BC anchor.** Per-stream MLP width 128, shared MLP $3 \times 128 \rightarrow 512 \rightarrow 512$. No recurrence. Actor heads are linear projections from the 512-dimensional embedding. Total: 776,190 parameters.
- **RL policy.** Per-stream MLP width 64, shared MLP $3 \times 64 \rightarrow 256 \rightarrow 256$. The 256-dimensional embedding is passed through a single-layer LSTM with input size 256 and hidden size 256 (PufferLib LSTMWrapper). Actor and critic heads are linear projections from the 256-dimensional LSTM output. Total: **TODO**.

Road segment features include a categorical type field that is replaced by a 7-class one-hot vector before encoding, expanding the road feature dimension from d_{road} to $d_{\text{road}} + 6$.

D Evaluation

D.1 Filtering the Waymo Dataset for Interactive SDC Scenarios

Most scenarios in the Waymo Open Motion Dataset (WOMD) involve the self-driving car (SDC) traveling without meaningful interaction with other agents—the SDC reaches its destination without requiring coordination or yielding. To increase the signal in our human-replay evaluation, we filter the dataset for scenarios in which the SDC trajectory intersects with other agents’ trajectories, indicating situations that require coordination, such as merging, yielding, or navigating busy intersections.

We score each scenario by counting the number of segment-level intersections between the SDC trajectory and all other agent trajectories, optionally filtering crossings that meet a minimum acute-

angle threshold (to exclude near-parallel overlaps, such as lane changes). From a pool of 10,000 held-out validation scenarios, we rank by intersection count and select the top 200 most interactive scenes. Figure 15 shows the resulting intersection count distributions across the full dataset and the selected subset, and Figure 16 shows nine representative examples from the selected set.




Figure 15: Distribution of SDC trajectory intersection counts. **Left:** raw intersection counts across all 50k scenarios. **Center:** angled intersections (non-zero only). **Right:** distribution within the selected top-200 subset.




Figure 16: Nine example scenarios from the selected interactive subset. The SDC trajectory is shown in green, other agents in blue, and trajectory intersection points with other logs in red.

D.2 Metrics

We report the following metrics across all experiments. Unless noted, all metrics are computed per active agent per episode and averaged across agents and scenarios.

Score. The primary metric. An agent scores 1 if it reaches its goal without any collision or off-road event during the episode, and 0 otherwise. It jointly captures all failure modes and is the main basis for comparing conditions.

Completion rate. The fraction of agents that reach their goal position (within $\delta = 2$ meters) before episode end, regardless of whether a collision or off-road event occurred.

Collision rate. The fraction of episodes in which the agent is involved in at least one collision with another vehicle.

At-fault collision rate. A subset of the collision rate taken from NAVSIM [58]. A collision is attributed to an agent if (i) the other vehicle is in front of the agent at the time of impact, and (ii) the

agent’s velocity vector points toward the other vehicle. This filters out collisions in which the agent was rear-ended or struck laterally by an inattentive partner.

Rear collision rate. The complement of at-fault: a collision is flagged as a rear collision for agent i if the colliding partner satisfies the at-fault conditions relative to i . A high rear collision rate indicates that the agent is being struck from behind, which is a signal that it is not keeping pace with traffic flow.

Collision severity (Δv). Beyond the binary collision indicator, we measure the severity of each at-fault collision event using the change in velocity (Δv) imparted to the agent at impact. Following the impulse-momentum formulation used in [64], the Delta-V of agent i in a collision with partner j is

$$\Delta v_i = \frac{m_j}{m_i + m_j} (1 + e) (\vec{v}_j - \vec{v}_i) \cdot \hat{n}, \quad (19)$$

where $\hat{n}$ is the unit collision normal (taken as the vector from agent i ’s center to agent j ’s center at impact), $e = 0.1$ is the coefficient of restitution for vehicle-to-vehicle crashes, and the dot product is clipped at zero to ignore separating velocities. Masses are proxied from bounding-box footprint for vehicles (anchored at 1500 kg for a 4.5 m $\times$ 1.8 m reference sedan) and fixed for vulnerable road users (75 kg for pedestrians, 90 kg for cyclists). Δv is one of the strongest predictors of injury risk in vehicle-to-vehicle crashes [64] and lets us distinguish low-impact contacts (e.g. parking-lot taps) from high-energy collisions even when the binary collision rate is identical. We report two summaries: the mean Δv conditional on collision, and the fraction of collisions with $\Delta v < 1$ mph (0.447 m/s), the latter is chosen for direct comparability with Waymo’s reporting.

Off-road rate. The fraction of episodes in which the agent crosses a road edge boundary, detected by checking for intersection between the agent bounding box and any road edge polyline.

Route progress ratio. Following [65], we measure how far along its expert reference trajectory each agent travels. At each timestep t , we find the closest point $x(t)$ on the agent’s logged trajectory and compute its arc-length distance $d_{x(t)}$ from the start of the path. The route progress ratio is

$$\rho = \frac{d_{x(t)} - d_p}{d_q - d_p}, \quad (20)$$

where d_p and d_q are the arc-length distances to the initial and final positions of the logged trajectory, respectively. A value of $\rho = 1$ means the agent reached its destination; $\rho > 1$ is possible if the agent overshoots. For agents that reach their goal under GOAL_REMOVE termination, we set $\rho = 1$ directly, since their position is invalidated upon removal. For all other agents, ρ is computed from the agent’s final position at episode end.

Lateral deviation. At each timestep t , we compute the Euclidean distance from the agent’s current position to the nearest point on its expert reference trajectory. We average this distance over all timesteps for which the agent is alive, yielding a mean lateral deviation in meters. This measures how far the agent drifts sideways from its intended route; a perfect replay would achieve zero lateral deviation.

ZSC gap. The difference in score between the self-play and human-replay evaluation conditions:

$$\Delta_{\text{ZSC}} = \text{score}_{\text{self-play}} - \text{score}_{\text{human-replay}}. \quad (21)$$

A value of zero indicates that the policy transfers without degradation to the human population.

E Additional Results

@Daphne: TODO

E.1 Convergence and Performance on the Training Dataset

Every policy is trained for 5 billion steps. Below, we plot the performance on the full training dataset as a sanity check for convergence. **TODO: In the plot below, you can clearly see that the regularized policies are not fully converged yet. My intuition is that once this is fixed, the replay perf will be much better as well.**




Figure 17: Performance of the policies trained on 50k maps on the full train dataset.

E.2 Human-Replay Results Comparison

Table 11: Human-replay performance on the randomly sampled validation set vs. the interactive validation subset for all scaling checkpoints. Top-3 values per column are highlighted (best, 2nd, 3rd); best value additionally in bold.

Self-play maps (metadata)	Anchor data (human demonstrations)	HR — validation (random)						HR — interactive					
		Score ↑	Coll. (%) ↓	At-fault (%) ↓	Rear coll. (%) ↓	Route prog. ↑	Lateral L2 ↓	Score ↑	Coll. (%) ↓	At-fault (%) ↓	Rear coll. (%) ↓	Route prog. ↑	Lateral L2 ↓
10	0 (unreg.)	0.637 ± 0.007	23.3 ± 0.6	21.0 ± 0.6	12.9 ± 0.5	0.989 ± 0.001	3.34 ± 0.11	0.378 ± 0.007	47.1 ± 0.7	43.8 ± 0.7	33.3 ± 0.7	0.990 ± 0.001	5.85 ± 0.15
100	0 (unreg.)	0.784 ± 0.006	15.7 ± 0.5	13.4 ± 0.5	7.9 ± 0.4	0.992 ± 0.001	2.50 ± 0.10	0.620 ± 0.007	31.9 ± 0.7	26.7 ± 0.6	20.1 ± 0.6	0.989 ± 0.001	5.27 ± 0.15
1k	0 (unreg.)	0.925 ± 0.004	5.1 ± 0.3	3.5 ± 0.3	2.5 ± 0.2	0.997 ± 0.001	1.40 ± 0.06	0.848 ± 0.005	8.9 ± 0.4	4.1 ± 0.3	6.0 ± 0.3	0.982 ± 0.002	3.69 ± 0.14
10k	0 (unreg.)	0.963 ± 0.003	1.9 ± 0.2	1.1 ± 0.1	1.0 ± 0.1	0.997 ± 0.001	1.21 ± 0.05	0.900 ± 0.004	6.0 ± 0.3	2.7 ± 0.2	5.2 ± 0.3	0.994 ± 0.001	2.72 ± 0.09
50k	0 (unreg.)	0.962 ± 0.003	1.9 ± 0.2	1.1 ± 0.1	1.2 ± 0.2	0.999 ± 0.000	1.07 ± 0.04	0.900 ± 0.004	6.9 ± 0.4	2.8 ± 0.2	6.0 ± 0.3	0.998 ± 0.000	2.03 ± 0.06
10	10 minutes	0.716 ± 0.006	15.1 ± 0.5	11.6 ± 0.4	8.0 ± 0.4	0.964 ± 0.002	0.87 ± 0.02	0.513 ± 0.007	30.2 ± 0.6	21.0 ± 0.6	21.9 ± 0.6	0.960 ± 0.003	1.58 ± 0.02
10	30 minutes	0.729 ± 0.006	14.4 ± 0.5	10.8 ± 0.4	8.0 ± 0.4	0.980 ± 0.002	1.33 ± 0.04	0.477 ± 0.007	34.1 ± 0.7	28.9 ± 0.6	21.1 ± 0.6	0.965 ± 0.002	2.22 ± 0.04
10	3 hours	0.738 ± 0.006	14.1 ± 0.5	11.3 ± 0.4	6.2 ± 0.3	0.977 ± 0.002	1.33 ± 0.05	0.464 ± 0.007	27.8 ± 0.6	22.0 ± 0.6	16.7 ± 0.5	0.948 ± 0.002	3.42 ± 0.09
10	30 hours	0.709 ± 0.006	11.7 ± 0.4	9.0 ± 0.4	6.2 ± 0.3	0.981 ± 0.002	1.01 ± 0.02	0.501 ± 0.007	32.2 ± 0.7	26.6 ± 0.6	19.8 ± 0.6	0.979 ± 0.002	1.66 ± 0.02
100	10 minutes	0.867 ± 0.005	9.3 ± 0.4	7.4 ± 0.4	3.8 ± 0.3	0.997 ± 0.000	0.73 ± 0.01	0.753 ± 0.006	18.8 ± 0.5	13.6 ± 0.5	9.3 ± 0.4	0.989 ± 0.001	1.46 ± 0.02
100	30 minutes	0.879 ± 0.005	8.1 ± 0.4	6.2 ± 0.3	2.7 ± 0.2	0.996 ± 0.001	0.72 ± 0.01	0.772 ± 0.006	16.8 ± 0.5	13.7 ± 0.5	4.9 ± 0.3	0.984 ± 0.001	1.20 ± 0.01
100	3 hours	0.907 ± 0.004	6.0 ± 0.3	4.4 ± 0.3	2.2 ± 0.2	0.995 ± 0.001	0.71 ± 0.01	0.797 ± 0.006	14.3 ± 0.5	7.8 ± 0.4	11.1 ± 0.4	0.989 ± 0.001	1.47 ± 0.02
100	30 hours	0.906 ± 0.004	6.9 ± 0.4	4.8 ± 0.3	2.2 ± 0.2	0.994 ± 0.001	0.81 ± 0.01	0.808 ± 0.006	14.8 ± 0.5	9.9 ± 0.4	9.3 ± 0.4	0.984 ± 0.002	1.51 ± 0.02
1k	10 minutes	0.956 ± 0.003	2.4 ± 0.2	1.0 ± 0.1	1.8 ± 0.2	0.997 ± 0.000	0.68 ± 0.01	0.920 ± 0.004	5.9 ± 0.3	2.2 ± 0.2	3.9 ± 0.3	0.992 ± 0.001	1.32 ± 0.02
1k	30 minutes	0.962 ± 0.003	2.6 ± 0.2	0.7 ± 0.1	2.1 ± 0.2	0.997 ± 0.000	0.68 ± 0.01	0.920 ± 0.004	7.5 ± 0.4	1.0 ± 0.1	7.1 ± 0.4	1.000	1.20 ± 0.01
1k	3 hours	0.968 ± 0.002	1.8 ± 0.2	0.7 ± 0.1	1.3 ± 0.2	0.997 ± 0.001	0.68 ± 0.02	0.930 ± 0.004	4.3 ± 0.3	0.5 ± 0.1	3.9 ± 0.3	0.992 ± 0.001	1.31 ± 0.02
1k	30 hours	0.959 ± 0.003	2.4 ± 0.2	1.0 ± 0.1	1.8 ± 0.2	0.997 ± 0.001	0.68 ± 0.01	0.900 ± 0.004	5.8 ± 0.3	2.1 ± 0.2	4.1 ± 0.3	0.985 ± 0.002	1.28 ± 0.02
10k	10 minutes	0.953 ± 0.003	1.5 ± 0.2	0.4 ± 0.1	1.2 ± 0.2	0.995 ± 0.001	0.74 ± 0.02	0.964 ± 0.003	2.3 ± 0.2	0.0	2.3 ± 0.2	0.997 ± 0.000	1.30 ± 0.02
10k	30 minutes	0.963 ± 0.003	1.5 ± 0.2	0.3 ± 0.1	1.3 ± 0.2	0.992 ± 0.001	0.67 ± 0.01	0.952 ± 0.003	3.5 ± 0.3	1.4 ± 0.2	3.5 ± 0.3	0.992 ± 0.001	1.19 ± 0.02
10k	3 hours	0.977 ± 0.002	1.2 ± 0.2	0.3 ± 0.1	1.0 ± 0.1	0.995 ± 0.001	0.78 ± 0.01	0.945 ± 0.003	3.5 ± 0.3	0.6 ± 0.1	2.5 ± 0.2	0.985 ± 0.002	1.45 ± 0.02
10k	30 hours	0.977 ± 0.002	1.3 ± 0.2	0.3 ± 0.1	1.2 ± 0.2	0.996 ± 0.001	0.68 ± 0.01	0.925 ± 0.004	3.5 ± 0.3	0.0	3.5 ± 0.3	0.985 ± 0.002	1.17 ± 0.01
50k	10 minutes	0.936 ± 0.003	1.6 ± 0.2	0.4 ± 0.1	1.3 ± 0.2	0.992 ± 0.001	0.77 ± 0.01	0.954 ± 0.003	1.1 ± 0.1	0.5 ± 0.1	1.1 ± 0.1	0.996 ± 0.000	1.41 ± 0.03
50k	30 minutes	0.974 ± 0.002	1.2 ± 0.1	0.2 ± 0.1	1.1 ± 0.1	0.994 ± 0.001	0.77 ± 0.01	0.936 ± 0.003	1.9 ± 0.3	1.5 ± 0.2	3.6 ± 0.3	0.991 ± 0.001	1.49 ± 0.02
50k	3 hours	0.980 ± 0.002	0.7 ± 0.1	0.2 ± 0.1	0.6 ± 0.1	0.997 ± 0.001	0.74 ± 0.01	0.927 ± 0.004	1.8 ± 0.2	0.0	1.8 ± 0.2	0.985 ± 0.001	1.32 ± 0.02
50k	30 hours	0.970 ± 0.002	1.1 ± 0.1	0.2 ± 0.1	1.0 ± 0.1	0.995 ± 0.001	0.70 ± 0.01	0.951 ± 0.003	1.9 ± 0.2	0.5 ± 0.1	1.5 ± 0.2	0.989 ± 0.001	1.24 ± 0.01

E.3 Ablations

We report the results on the following ablations.

Sensor noise.

Lambda conditioning.

Reward conditioning.

F Mapping Agent Experience to Human Time

We train self-play RL agents on 10 billion transitions. Given that the Waymo scenarios are discretized at 10 Hz, each transition (o_t, a_t) represents 0.1 seconds. The total training experience, therefore, corresponds to 31.68 years of real-time interaction.

SMART [42] was trained on the full Waymo training dataset, which has 500,000 training scenarios. At 10 Hz, this means 45 million transitions. We train checkpoints on the SDC trajectory from 67, 200, 1200 and 12,000 maps, which roughly has 90 transitions per scenario.